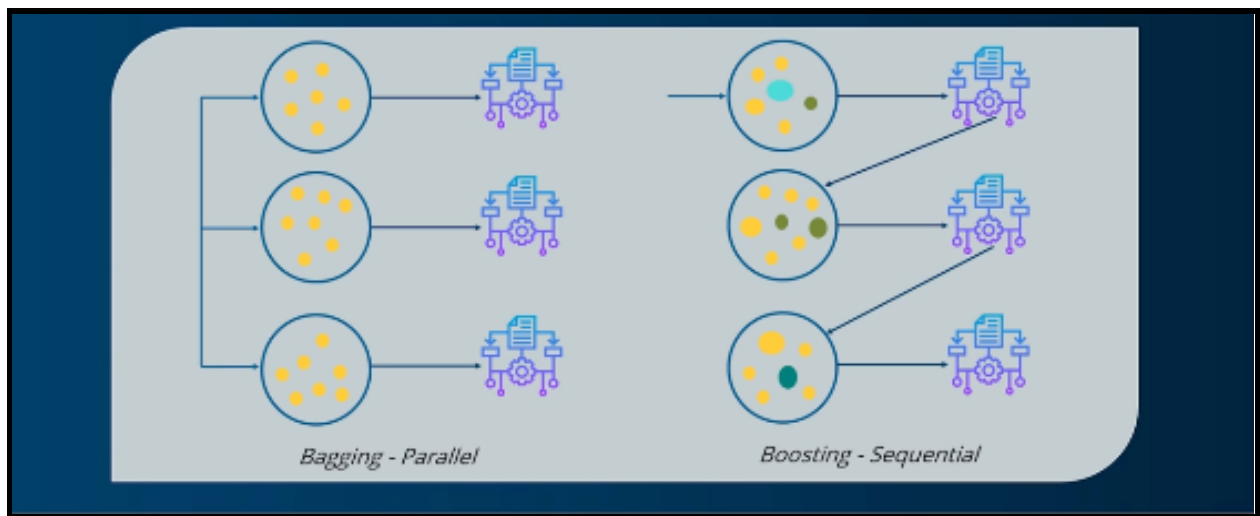


Machine Learning technique - Boosting

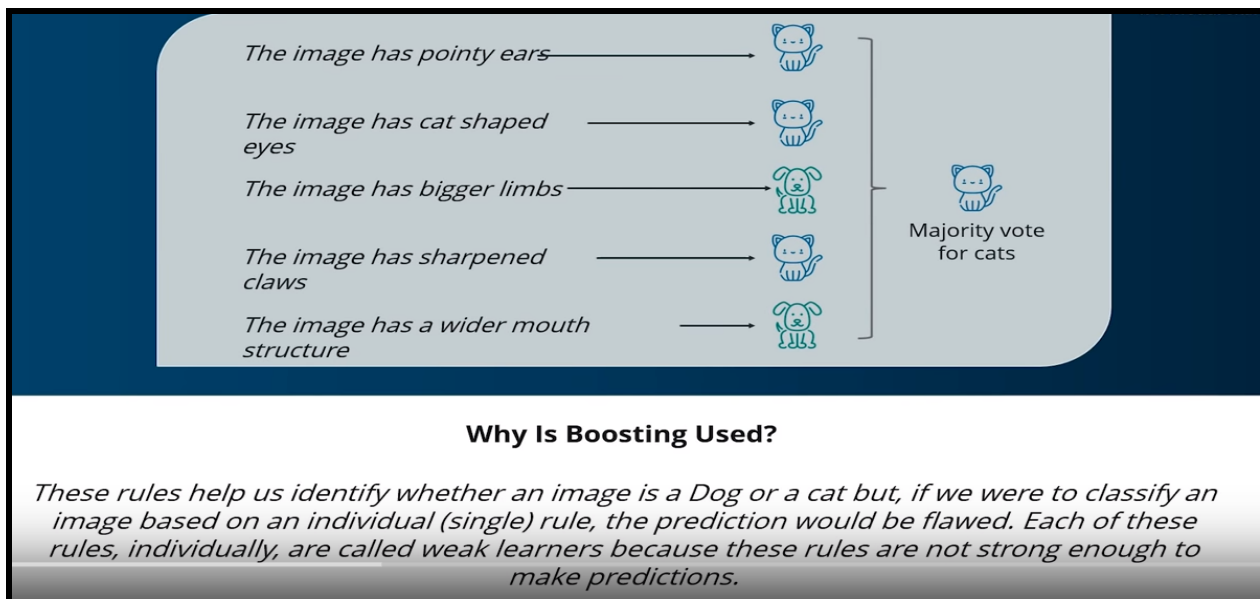
Boosting technique is ensemble learner which takes only one model either linear or decision tree, whereas by default it uses the decision tree.

What is Ensemble: Ensemble learning offers a systematic solution to combine the predictive power of multiple learners. The resultant is a single model which gives the aggregated output from several models.

Image: 1: Workflow diagram for bagging and boosting ensemble




Why we need a Boosting technique - Let's take the example of classifying cat and dog image.



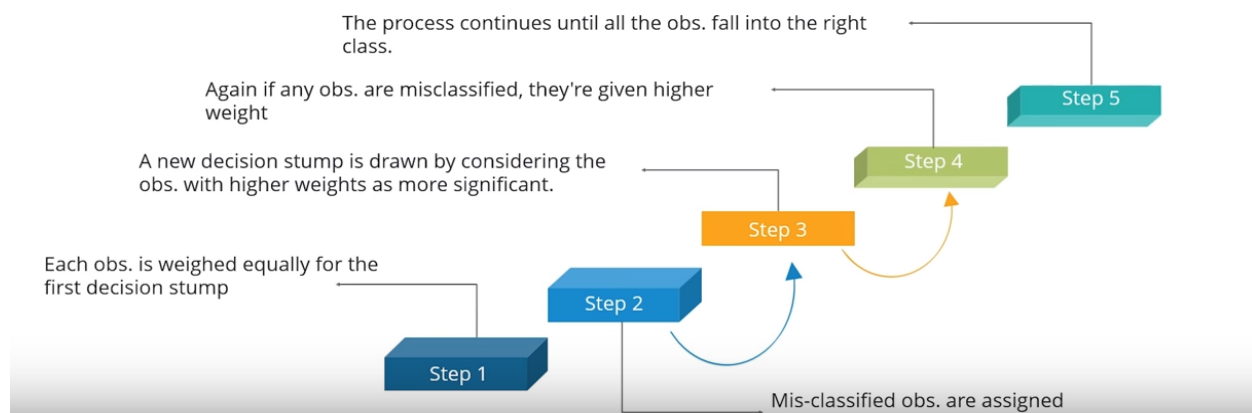
WHAT IS BOOSTING?

Boosting is a process that uses a set of Machine Learning algorithms to combine weak learner to form strong learners in order to increase the accuracy of the model.



If you see Image1: It's showing the way boosting method works. Every time we pass the complete data set to a weak learner (model) and for whichever records it's not able to predict the output correctly we give more weightage to that particular record so that in our next weak learner that particular records get predicted perfectly.

ADAPTIVE BOOSTING

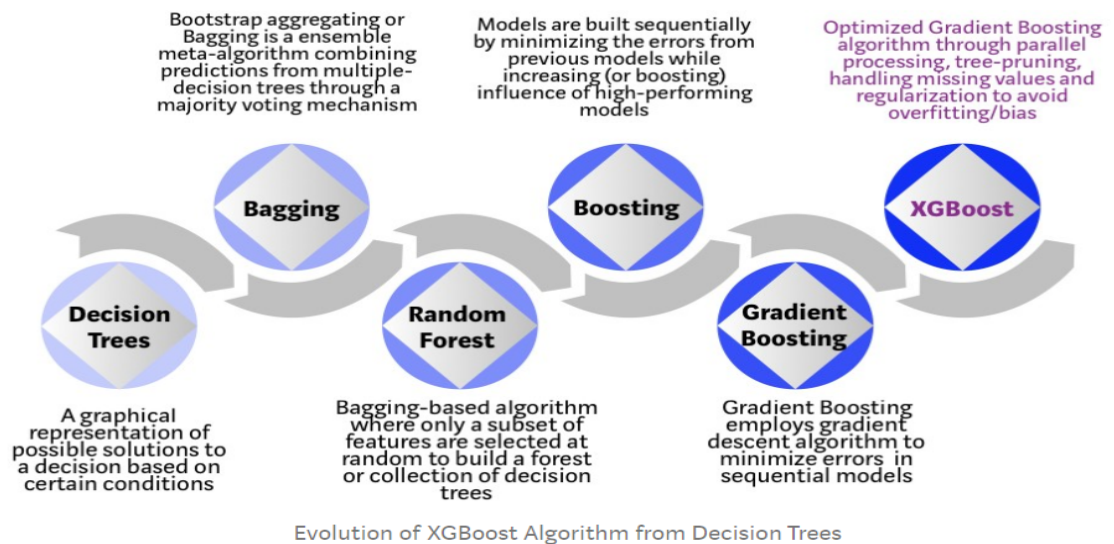


Step2: Misclassified obs. are assigned higher weightage

Gradient Boosting: It is almost similar to Adaboost but here we don't do weight updation for misclassified records. Instead of weight updation, we optimize the loss function of the previous model. (takes longer time to compute)

- Components used-
1. Loss function
 2. Weak learner (model)
 3. The additive model that regularize the loss function (new model)

XG BOOST: It's the updated version of gradient boosting which focuses on computational power and model efficiency.



Key features -

- Parallelization (create many trees parallely)
- Distributed computing (helps to evaluate quickly)
- Out-of Core computing (able to handle varied data set)
- Cache optimization (enables to use best possible of your hardware configuration)



Maths behind GBM or XG Boost:-

Let's understand it using game

Let's play a game...

You are given $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$, and the task is to fit a model $F(x)$ to minimize square loss.

Suppose your friend wants to help you and gives you a model F .

You check his model and find the model is good but not perfect.

There are some mistakes: $F(x_1) = 0.8$, while $y_1 = 0.9$, and

$F(x_2) = 1.4$ while $y_2 = 1.3$... How can you improve this model?

Rule of the game:

- ▶ You are not allowed to remove anything from F or change any parameter in F .
- ▶ You can add an additional model (regression tree) h to F , so the new prediction will be $F(x) + h(x)$.

Simple solution:

You wish to improve the model such that

$$F(x_1) + h(x_1) = y_1$$

$$F(x_2) + h(x_2) = y_2$$

...

$$F(x_n) + h(x_n) = y_n$$

Or, equivalently, you wish

$$h(x_1) = y_1 - F(x_1)$$

$$h(x_2) = y_2 - F(x_2)$$

...

$$h(x_n) = y_n - F(x_n)$$

Just fit a regression tree h to data

$$(x_1, y_1 - F(x_1)), (x_2, y_2 - F(x_2)), \dots, (x_n, y_n - F(x_n))$$

$y_i - F(x_i)$ are called **residuals**. These are the parts that existing model F cannot do well.

The role of h is to compensate the shortcoming of existing model F .




If the new model $F + h$ is still not satisfactory, we can add another regression tree...

How is this related to gradient descent?

Gradient Descent :

Minimize a function by moving in the opposite direction of the gradient.

$$\theta_i := \theta_i - \rho \frac{\partial J}{\partial \theta_i}$$



Loss function $L(y, F(x)) = (y - F(x))^2/2$

We want to minimize $J = \sum_i L(y_i, F(x_i))$ by adjusting $F(x_1), F(x_2), \dots, F(x_n)$.

$$\frac{\partial J}{\partial F(x_i)} = \frac{\partial \sum_i L(y_i, F(x_i))}{\partial F(x_i)} = \frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} = F(x_i) - y_i$$

So we can interpret residuals as negative gradients.

$$y_i - F(x_i) = -\frac{\partial J}{\partial F(x_i)}$$

Relationship between Gradient descent and GBM:

$$F(x_i) := F(x_i) + h(x_i)$$

$$F(x_i) := F(x_i) + y_i - F(x_i)$$

$$F(x_i) := F(x_i) - 1 \frac{\partial J}{\partial F(x_i)}$$

$$\theta_i := \theta_i - \rho \frac{\partial J}{\partial \theta_i}$$

For regression with **square loss**,

$$\text{residual} \Leftrightarrow \text{negative gradient}$$

$$\text{fit } h \text{ to residual} \Leftrightarrow \text{fit } h \text{ to negative gradient}$$

$$\text{update } F \text{ based on residual} \Leftrightarrow \text{update } F \text{ based on negative gradient}$$

So we are actually updating our model using **gradient descent**!

It turns out that the concept of **gradients** is more general and useful than the concept of **residuals**.